

# Apache Kafka, what is it good for?

Leotis Buchanan

May 5, 2017

# Why I luv Kafka

This was linkedin architecture before kafka

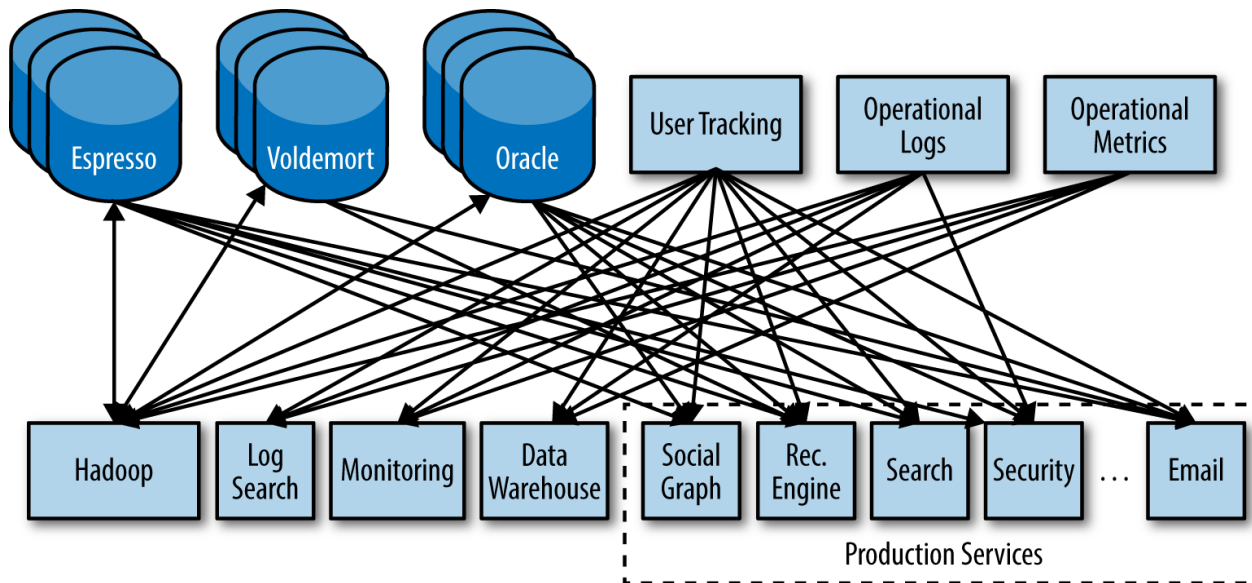


image borrowed from: "Jay Kreps why I love logs book"

# why I luv Kafka

look what happened after they decided to use kafka

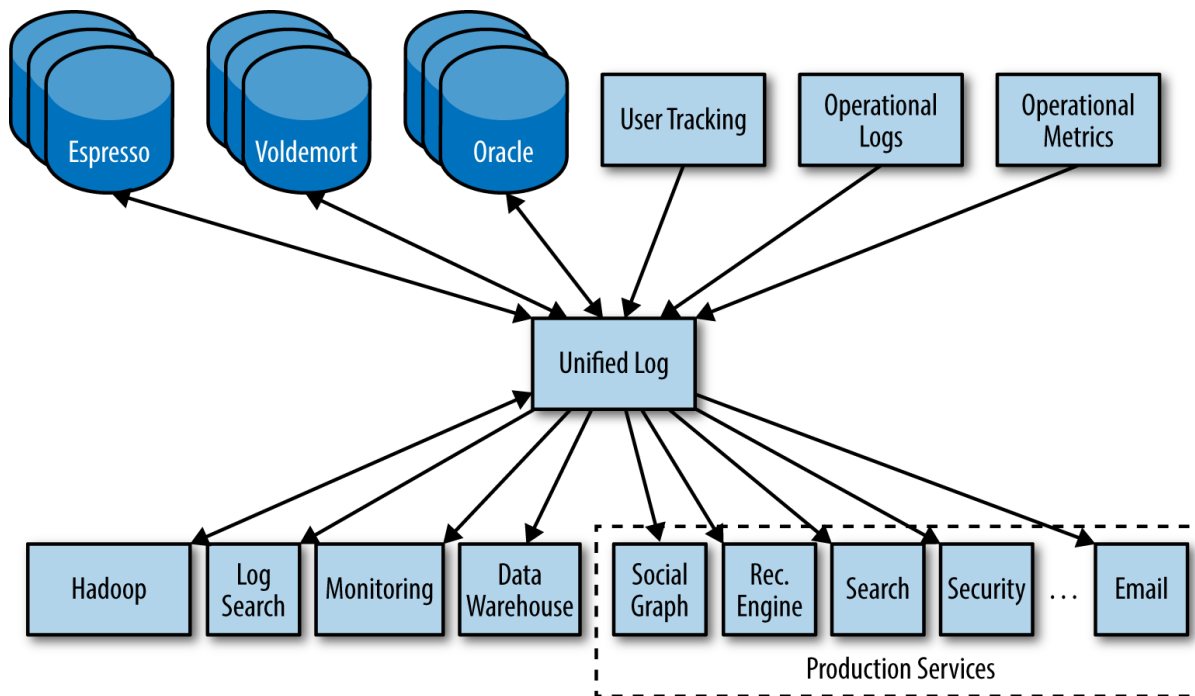


image borrowed from: "Jay Kreps why I love logs book"

# why I luv Kafka

**They have been able to[1]:**

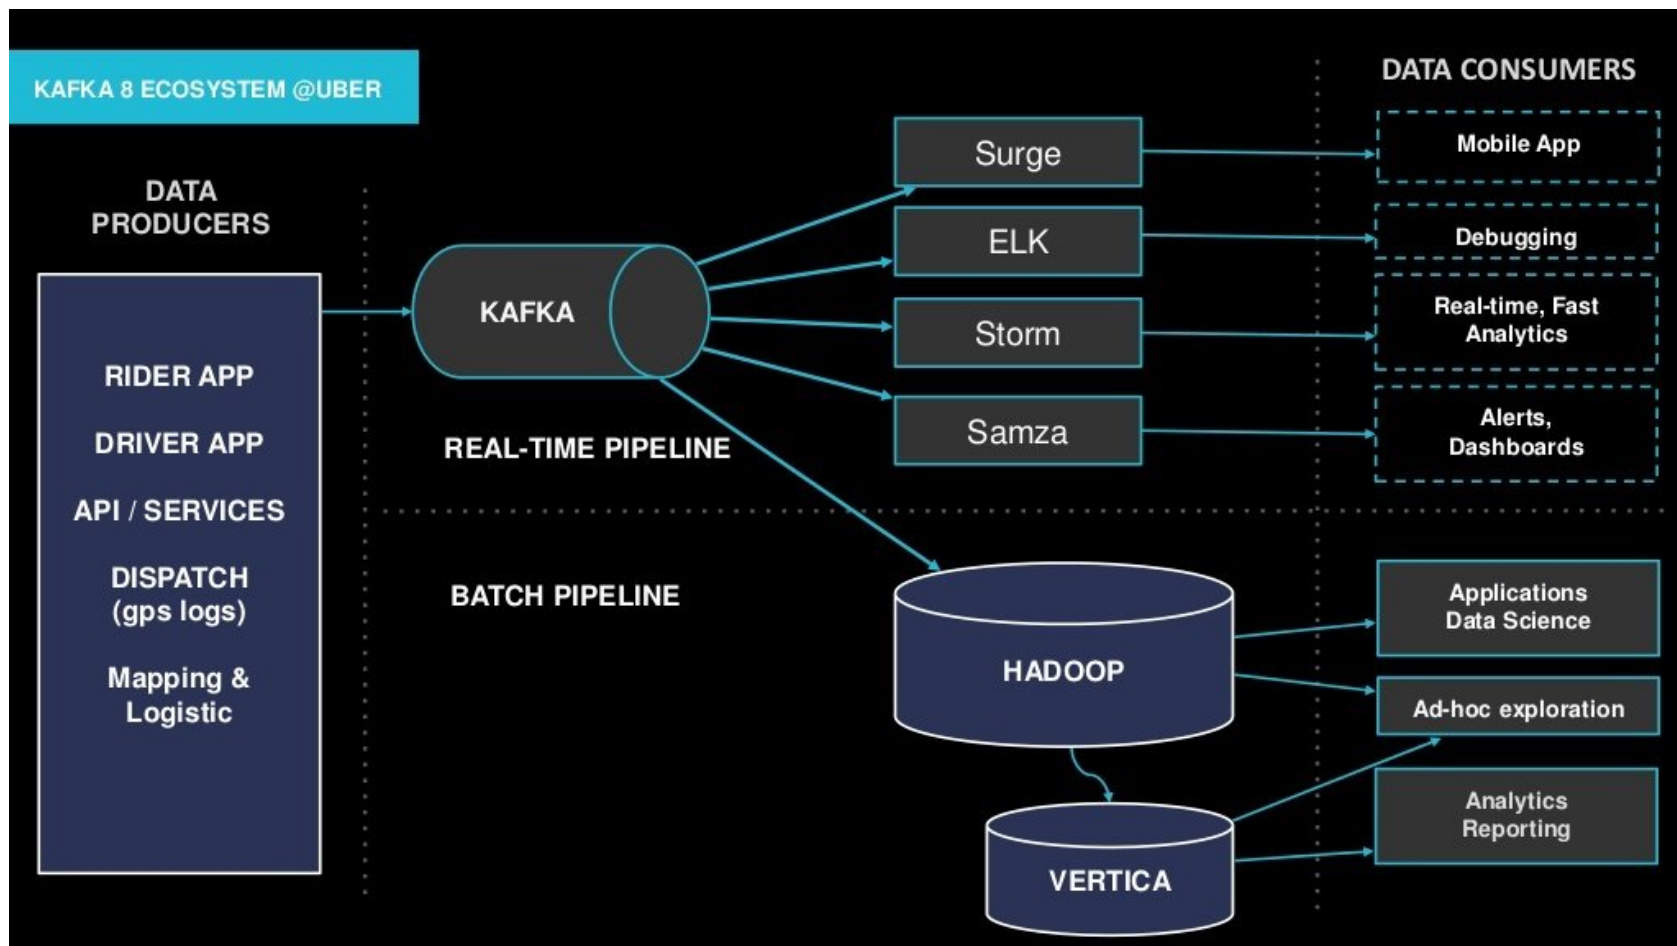
- capture user activity in real time
- change from a batch oriented system to real-time data processing system
- handle 10 billion message writes each day with a sustained peak of over 172,000 messages per second
- support dozens of subscribing systems and delivers more than 55 billion messages to these consumer processing each day

# why I luv Kafka

## **Who else is using kafka to do some really cool things**

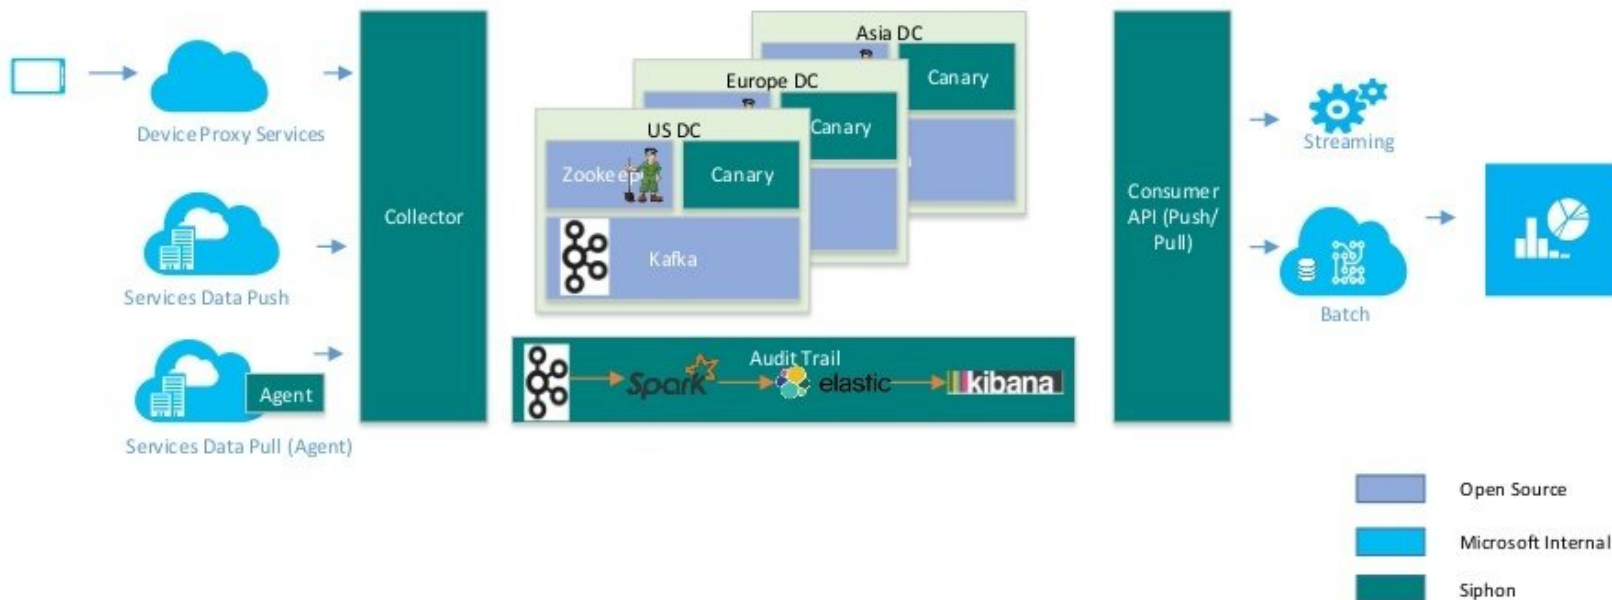
- uber is using it to build something they all "The World's Realtime Transit Infrastructure"[2]

## Uber "The World's Realtime Transit Infrastructure"[2]



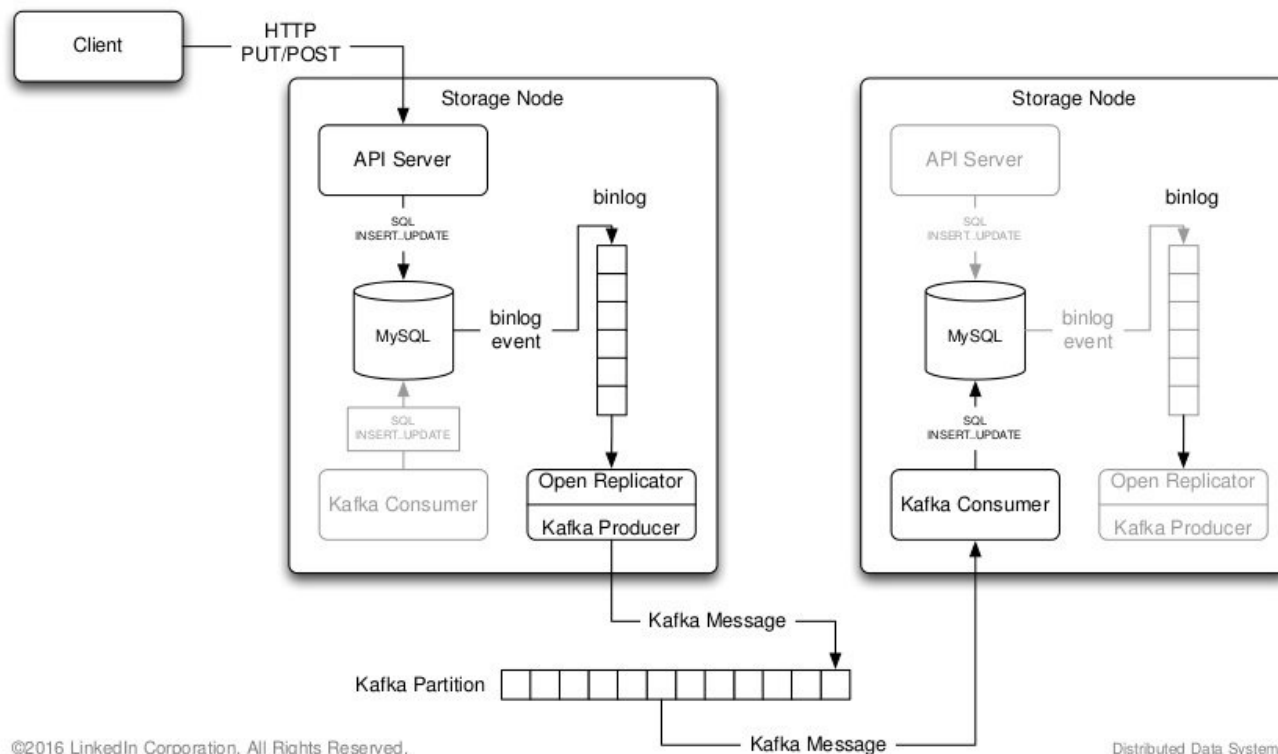
## Microsoft "Siphon"[3]

# Siphon Architecture



## Linkedin "espresso database replication with kafka"[4]

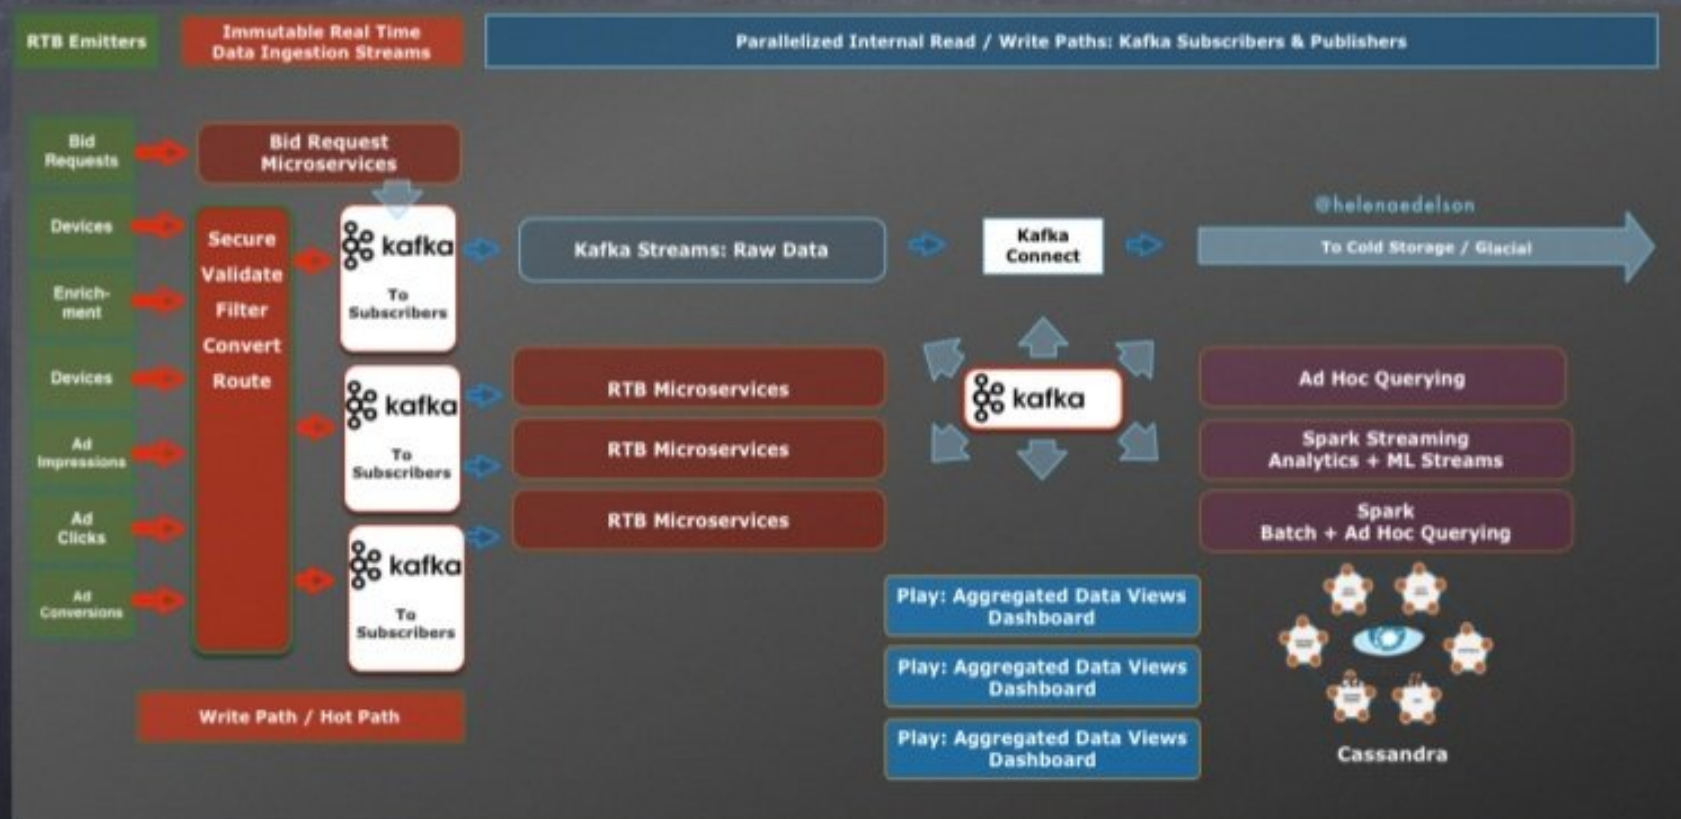
### Kafka for Internal Replication





**TupleJump "Leveraging Kafka for Big Data in Real Time Bidding, Analytics, Machine Learning and Campaign Management for Globally Distributed Data Flows"[5]**

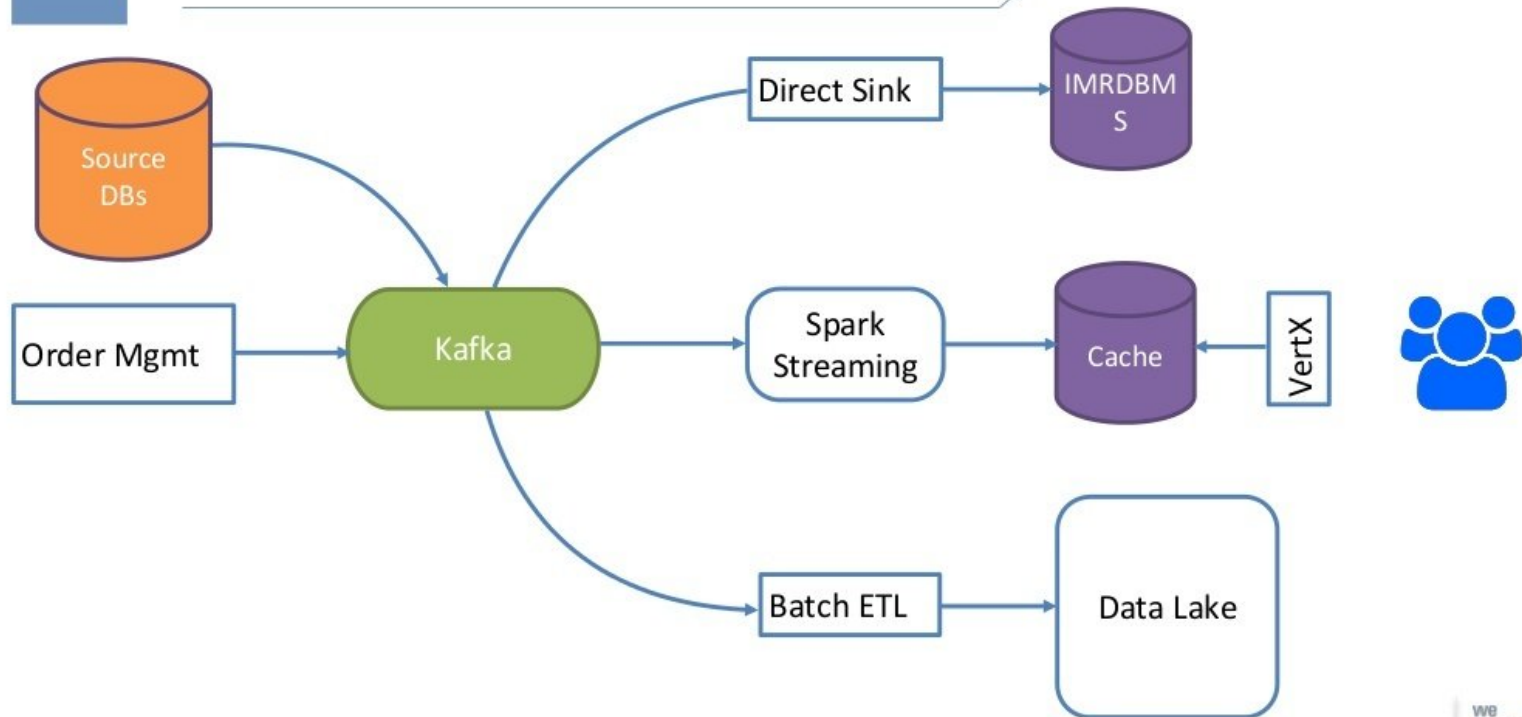
# Kafka as Platform Fabric



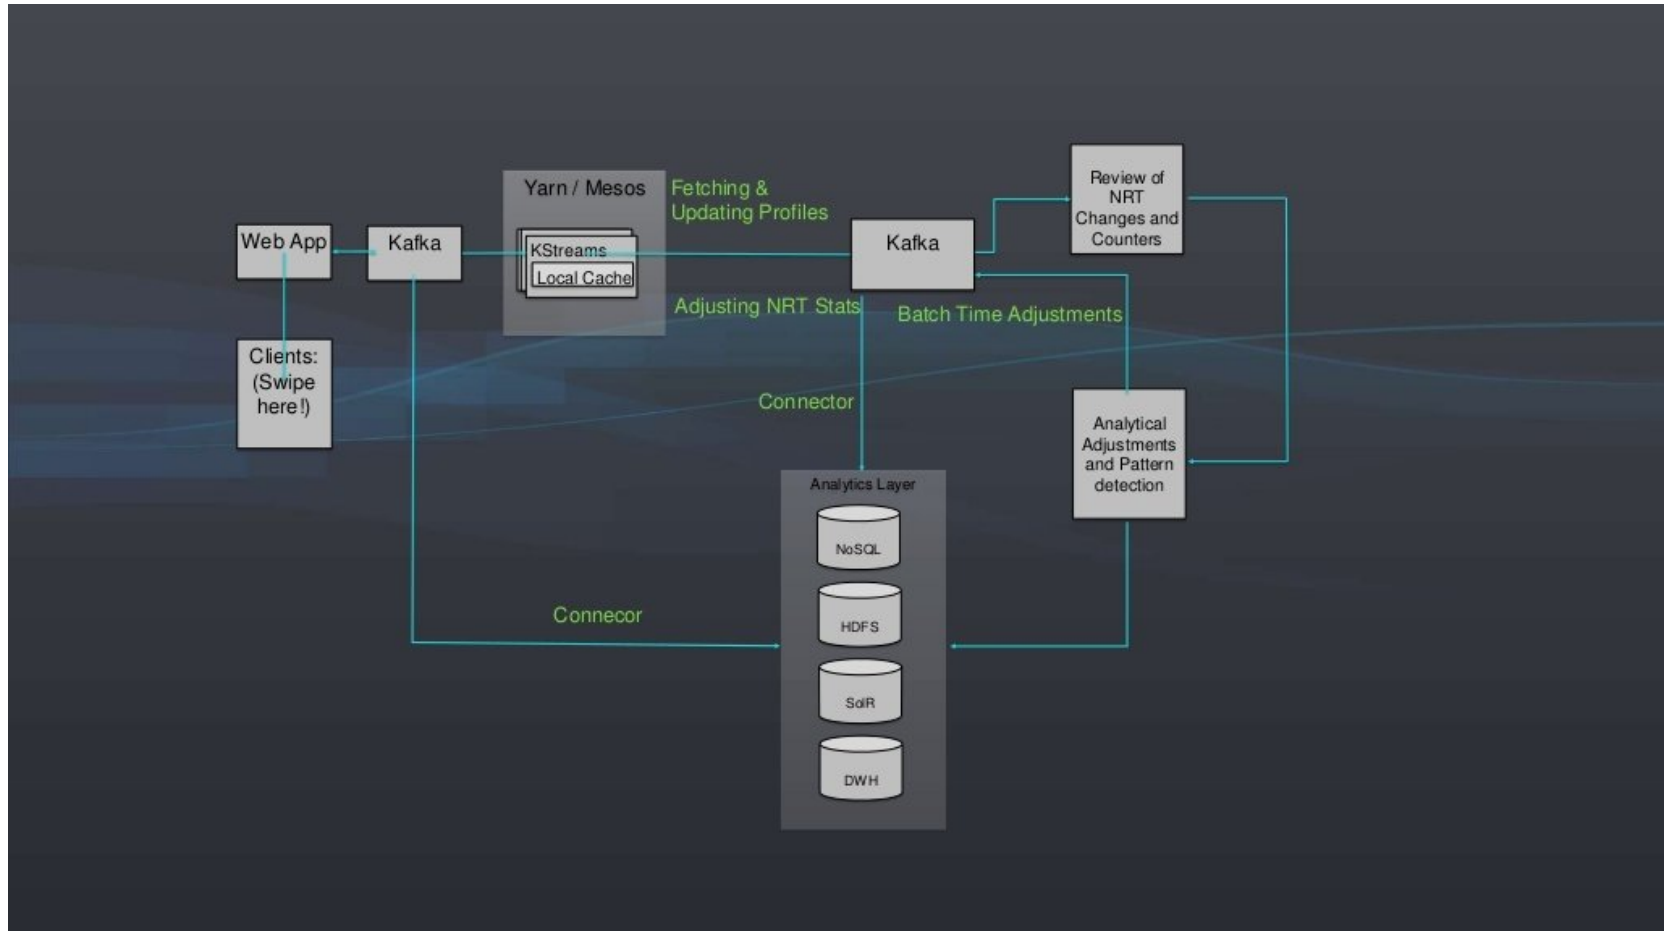
## Goldman Sachs "Real-Time Analytics Visualized w/ Kafka + Streamliner + MemSQL + ZoomData" [6]



### Representative Deployment



## Fraud detection "Gwen Shapira" [7]




# What is kafka?

- Originally created at LinkedIn in 2010
- Designed to support batch and real-time analytics
- Performs extremely well at very large scale
- LinkedIn's installation of Kafka processes over 1.4 *trillion* messages per day
- Made open source in 2011, became a top-level Apache project in 2012
- In use at many not very successful organizations:
- Twitter, Netflix, Goldman Sachs, Hotels.com, IBM, Spotify, Uber, Square, Cisco...

# Motivation for kafka?

You start out with a simple requirement, so you decide to build a simple pipeline


- A single place where all data resides
- A single ETL (Extract, Transform, Load) process to move data to that location



# Motivation for kafka?

It is a beautiful success, now everybody knows about it and now they need you add more features:

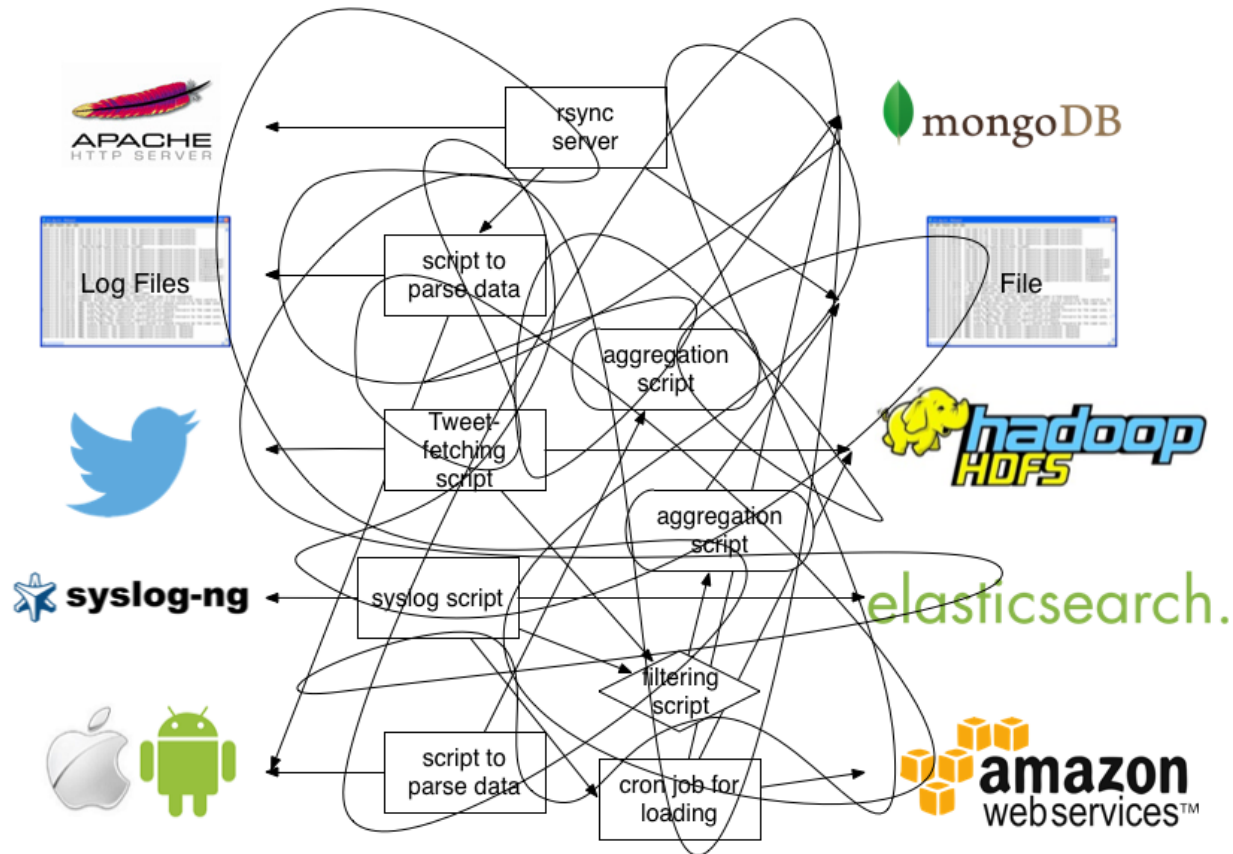
1. every body wants to be notified if it fails
2. can we add more datasources?
3. they have a new vision for your hastily put together pipeline?
4. Lets make this available to everyone!!





**Yes they are serious. So  
you reluctantly follow  
orders**

# and create a master piece



Then it fails a month after and they ask  
you to fix it



# so back to the motivations for kafka

- Data pipelines typically start out simply
- Data pipelines inevitably grow over time
- Systems and ETL become increasingly hard to manage

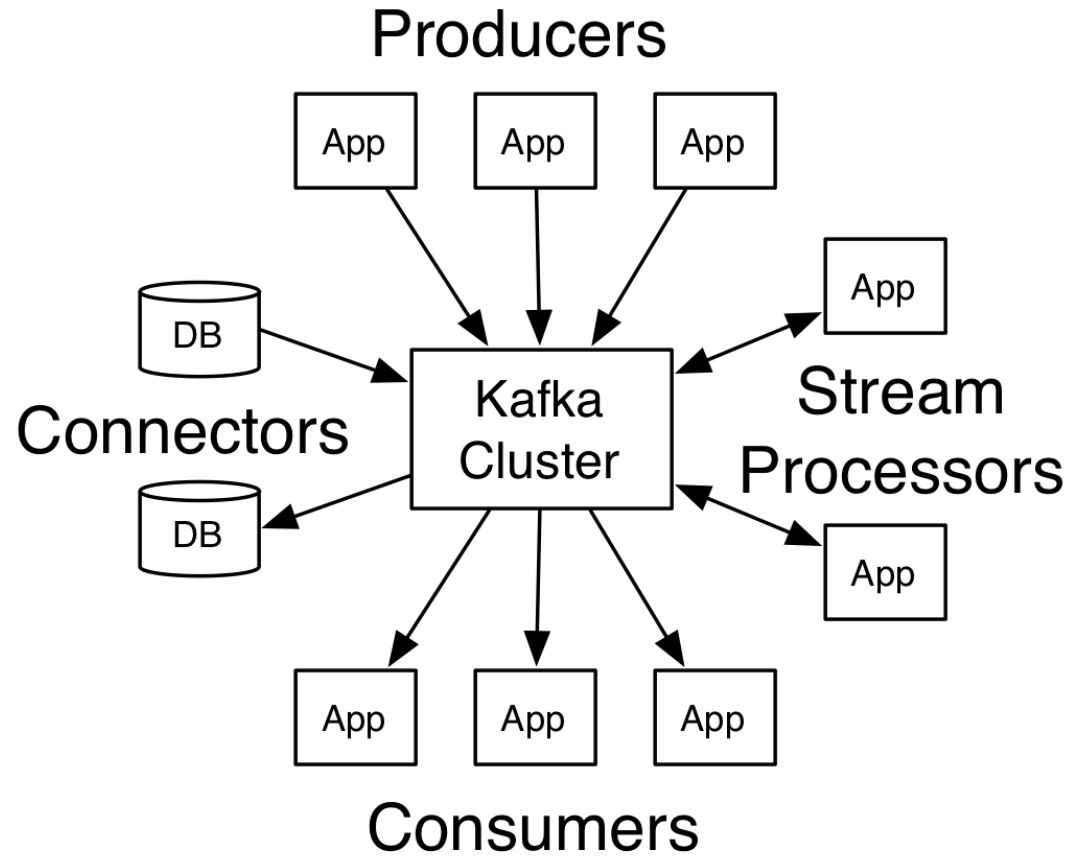
# so back to the motivations for kafka

- Traditionally, almost all data processing was batch-oriented
- This is limiting
- If you process data as it is being generated you can make decisions now.

# Kafka is a universal pipeline

- Kafka decouples data source and destination systems – Via a *publish/subscribe* architecture
- Using Kafka, all data sources write their data to the Kafka cluster
- Any system wishing to use the data reads it from Kafka

# Kafka is a universal pipeline



# Kafka is a universal pipeline

- Once the data is in Kafka, it can be read by multiple different Consumers
- Increasing the number of Consumers does not add significant load to the system
- Adding a new Consumer does not require any modification to the Producer(s)



# Kafka features

- Producers write data in the form of *messages* to the Kafka cluster
- Messages are written to *topics*
- Consumers read messages from one or more topics
- Data retention time in Kafka can be configured on a per-topic basis

# Kafka features

- **Kafka is very scalable, and very resilient**
  - Even a small cluster can process a large volume of messages
  - Tests have shown that three low-end machines can easily deal with two million writes per second
  - Messages are replicated on multiple machines for reliability
- **Consumers can be shut down temporarily**
  - When they restart, they will continue to read from where they left off

# Kafka Components

**There are four key components in a Kafka system**

- Producers
- Brokers
- Consumers
- ZooKeeper

# Producers

- A Producer sends messages to the Kafka cluster
- Producers can be written in any language
- A command-line Producer tool exists to send messages to the cluster

# Kafka messages

- A message is the basic unit of data in Kafka
- A message is a key-value pair
- Key and value can be any data type
  - You provide a serializer to turn the key and value into byte arrays
- Key is optional
  - Keys are used to determine which *Partition* (see later) a message will be sent to

# Topics

- Each message belongs to a *Topic*
- Developers decide which topics exist
- Typically, different systems will write to different topics

# Broker

- Brokers receive and store messages when they are sent by the Producers
- A Kafka cluster will typically have multiple Brokers
- Each Broker manages one or more Partitions

# Brokers Manage Partitions

- Any given Partition is handled by a single Broker
- Each Partition is stored on the Broker's disk as one or more log files
- Each message in the log is identified by its *offset*



# Consumer Basics

- Consumers pull messages from the cluster
- Multiple Consumers can read data from the same topic

# Why Pull messages

- Kafka consumers work by pulling messages
  - The advantages of pulling, rather than pushing, data, include:
    - The ability to add more Consumers to the system without reconfiguring the cluster
    - The ability for a Consumer to go offline and return later, resuming from where it left off
    - No problems with the Consumer being overwhelmed by data
    - It can pull, and process, the data at whatever speed it needs to
-

# Keeping Track of Position

- As messages are written to a topic, the Consumer will automatically retrieve them
- The *Consumer Offset* keeps track of the latest message read
- If necessary, the Consumer Offset can be changed
- The Consumer Offset is stored in a special Kafka topic

# ZooKeeper

- Apache ZooKeeper is an Apache project
- It is “a centralized service for maintaining configuration information”
- Used by many projects

# How Kafka Uses ZooKeeper

Kafka Brokers use ZooKeeper for a number of important internal features

- Leader election, failure detection

# Decoupling Producers and Consumers

- A key feature of Kafka is that Producers and Consumers are decoupled
- A slow Consumer will not affect Producers
- More Consumers can be added without affecting Producers
- Failure of a Consumer will not affect the system
- Multiple brokers, multiple topics, and *Consumer Groups* provide very high scalability

**Example code**

# Creating a producer

```
private Properties kafkaProps = new Properties();
kafkaProps.put("bootstrap.servers",
    "broker1:9092,broker2:9092");

kafkaProps.put("key.serializer",
    "org.apache.kafka.common.serialization.String-
        Serializer");

kafkaProps.put("value.serializer",
    "org.apache.kafka.common.serializa-
        tion.StringSerializer");

producer = new KafkaProducer<String, String>(kafkaProps);
```



# Sending a message

```
ProducerRecord<String, String> record =  
    new ProducerRecord<>( "CustomerCountry",  
        "Precision Products",  
        "France" );  
try {  
    producer.send(record);  
} catch (Exception e) {  
}
```

# Creating a consumer

```
Properties props = new Properties();

props.put("bootstrap.servers",
"broker1:9092,broker2:9092");
props.put("group.id", "CountryCounter");
props.put("key.deserializer",
"org.apache.kafka.common.serialization.StringDe-
serializer");
props.put("value.deserializer",
"org.apache.kafka.common.serialization.StringDe-
serializer");
KafkaConsumer<String, String> consumer =
new KafkaConsumer<String,
String>(props);
```

# Subscribing to a topic

```
consumer.subscribe(  
Collections.singletonList("customerCountries"));
```

# Reading a message

```
try {
    while (true) {
        ConsumerRecords<String, String> records = consumer.poll(100);
        for (ConsumerRecord<String, String> record : records)
        {
            log.debug("topic = %s, partition = %s, offset = %d, customer = %s,
country = %s\n",
record.topic(), record.partition(), record.offset(), record.key(),
record.value());
            int updatedCount = 1;
            if (custCountryMap.containsKey(record.value())) {
                updatedCount = custCountryMap.get(record.value()) + 1;
            }
            custCountryMap.put(record.value(), updatedCount)
            JSONObject json = new JSONObject(custCountryMap);
            System.out.println(json.toString(4))
        }
    }
} finally {
    consumer.close();
}
```

# Summary

This was an introductory presentation. Kafka is a really cool tool.

if you want to learn more you can:

- Do the confluent kafka course  
[<https://www.confluent.io/training/>]
- Get yourself a copy of the kafka definite guide book  
[<http://shop.oreilly.com/product/0636920044123.do>]
- Build something using kafka.

**Thank you for listening**

**Easy Questions Please**

# . References

- 1. Building LinkedIn's Real-time Activity Data Pipeline Ken Goodhope, Joel Koshy, Jay Kreps, Neha Narkhede, Richard Park, Jun Rao, Victor Yang Ye LinkedIn
- 2. Kafka + Uber- The World's Realtime Transit Infrastructure, Aaron Schildkrout [<https://www.confluent.io/kafka-summit-2016-keynote-kafka-and-uber-the-worlds-realtime-transit-infrastructure/>]
- 3. Siphon: [<https://www.confluent.io/kafka-summit-2016-users-siphon-near-real-time-databus-using-kafka/>]
- 4. Espresso Database Replication with Kafka [<https://www.confluent.io/kafka-summit-2016-users-espresso-database-replication-with-kafka/>]
- 5. Goldma sacs [<https://www.confluent.io/kafka-summit-2016-users-real-time-analytics-visualized-with-kafka/>]